

CSCB63 – Design and Analysis of Data Structures

Akshay Arun Bapat¹

¹based on notes by Anya Tafliovich, Anna Bretscher and Albert Lai

Mergeable heaps

Recall the heap data structure:


- `insert(j , p)`: insert job j with priority p
- `max()` or `min()`: return job with max/min priority
- `extract-max()` or `extract-min()`: remove and return job with max/min priority
- `change-priority(j , p')`: increase or decrease priority of job j to p' (optional)

Does not support:

- `union(H_1 , H_2)`: merge / union two heaps H_1 and H_2

Fibonacci (min-)heap

- a forest of (min-)heaps:
 - parent priority \leq child priority
 - siblings in circular doubly-linked list; parent points to one arbitrary child
- roots in circular doubly-linked list
- pointer to minimum-priority root



Binary heap vs Fibonacci heap

	binary heap worst-case	Fibonacci heap amortised
insert	$\Theta(\log n)$	$\Theta(1)$
extract-min	$\Theta(\log n)$	$O(\log n)$
decrease-priority	$\Theta(\log n)$	$\Theta(1)$
union	$\Theta(n)$	$\Theta(1)$

If Prim's algorithm uses a Fibonacci heap:

- if $n = |V|$ and $m = |E|$, then we have
- n calls of extract-min: $O(n \log n)$ total
- and up to m calls of decrease-priority: $O(m)$ total

for a total of: $O(n \log n + m)$ time

Fibonacci heap: fields

Each node has:

- *key*: priority
- *left*, *right*: for circular list of siblings
- *parent*: pointer to parent
- *child*: pointer to one child
- *degree*: number of children
- *marked*: boolean, important during decrease-priority

The heap has:

- *root_list*: a circular doubly-linked list of roots of the heaps
- *min*: pointer to root node with minimum *key*

Fibonacci heap: insert


`insert(H, k):`

0. `new_root := new node(key=k, marked=false)`

1. add `new_node` to `H.root_list`

2. if `k < H.min.key`:

3. `H.min = new_root`



Fibonacci heap: insert


`insert(H, k):`

0. `new_root := new node(key=k, marked=false)`

1. add `new_node` to `H.root_list`

2. if `k < H.min.key`:

3. `H.min = new_root`



Fibonacci heap: insert


`insert(H, k):`

0. `new_root := new node(key=k, marked=false)`

1. add `new_node` to `H.root_list`

2. if `k < H.min.key`:


3. `H.min = new_root`



Fibonacci heap: union

`union(H, H_1, H_2):`


0. `H.root_list := H_1.root_list + H_2.root_list`
1. `if H_1.min.key <= H_2.min.key:`
2. `H.min := H_1.min`
3. `else:`
4. `H.min := H_2.min`



Fibonacci heap: union

`union(H, H_1, H_2):`


0. `H.root_list := H_1.root_list + H_2.root_list`
1. `if H_1.min.key <= H_2.min.key:`
2. `H.min := H_1.min`
3. `else:`
4. `H.min := H_2.min`



Fibonacci heap: union

`union(H, H_1, H_2):`

0. `H.root_list := H_1.root_list + H_2.root_list`
1. `if H_1.min.key <= H_2.min.key:`
2. `H.min := H_1.min`
3. `else:`
4. `H.min := H_2.min`



Fibonacci heap: insert and union

- Complexity of insert: $\mathcal{O}(1)$
- Complexity of union: $\mathcal{O}(1)$
- “Real work” is in extract-min and decrease-priority

Fibonacci heap: extract-min


extract-min(H):

0. remove H.min from H.root_list

1. add each child of H.min to H.root_list

2. H.min := any former child of H.min // can be wrong!

3. consolidate(H) // real work here



Fibonacci heap: extract-min


`extract-min(H):`

0. remove `H.min` from `H.root_list`

1. add each child of `H.min` to `H.root_list`

2. `H.min := any former child of H.min` // can be wrong!

3. `consolidate(H)` // real work here



Fibonacci heap: extract-min


`extract-min(H):`

0. remove `H.min` from `H.root_list`

1. add each child of `H.min` to `H.root_list`

2. `H.min := any former child of H.min` // can be wrong!

3. `consolidate(H)` // real work here



consolidate: idea

Want:

- end with root list with nodes of unique degree


Idea:

- repeat until all nodes in root list have unique degree:
 - walk through root list
 - remember degree of each node so far
 - if see a node x with degree same as that of already seen y ,
 - $u := x$ or y , whoever's key is larger
 - $v := x$ or y , whoever's key is smaller
 - add u to children of v
 - remove u from root list
- update *min*

How to remember degrees of nodes?

- maintain array A of pointers
- $A[i]$ is root node with degree i

consolidate: example




$A[0] = nil$

$A[1] = nil$

$A[2] = nil$

$A[3] = nil$

consolidate: example




$A[0] = \text{nil}$

$A[1] = 18$

$A[2] = \text{nil}$

$A[3] = \text{nil}$

consolidate: example




$A[0] = \text{nil}$

$A[1] = 18$

$A[2] = \text{nil}$

$A[3] = \text{nil}$

consolidate: example




$A[0] = 52$

$A[1] = 18$

$A[2] = \text{nil}$

$A[3] = \text{nil}$

consolidate: example




$A[0] = 52$

$A[1] = 18$

$A[2] = \text{nil}$

$A[3] = \text{nil}$

consolidate: example




$A[0] = 52$

$A[1] = 18$

$A[2] = \text{nil}$

$A[3] = \text{nil}$

consolidate: example




$A[0] =$ (52)

$A[1] = nil$

$A[2] =$ (18)

$A[3] = nil$

consolidate: example




$A[0] =$ (52)

$A[1] =$ (17)

$A[2] =$ (18)

$A[3] = nil$

consolidate: example




$A[0] =$ (52)

$A[1] =$ (17)

$A[2] =$ (18)

$A[3] = nil$

consolidate: example




$A[0] =$ (52)

$A[1] =$ (17)

$A[2] =$ (18)

$A[3] = nil$

consolidate: example




$A[0] =$ (52)

$A[1] =$ (17)

$A[2] = \text{nil}$

$A[3] =$ (18)

consolidate: example




$A[0] = (52)$

$A[1] = (17)$

$A[2] = \text{nil}$

$A[3] = (18)$

consolidate: example




$A[0] = nil$

$A[1] = (17)$

$A[2] = nil$

$A[3] = (18)$

consolidate: example




$A[0] = nil$

$A[1] = 17$

$A[2] = nil$

$A[3] = 18$

consolidate: example




$A[0] = \text{nil}$

$A[1] = \text{nil}$

$A[2] = 17$

$A[3] = 18$

consolidate: example




$A[0] = \textcircled{7}$

$A[1] = \textit{nil}$

$A[2] = \textcircled{17}$

$A[3] = \textcircled{18}$

consolidate: example




$A[0] = 7$

$A[1] = nil$

$A[2] = 17$

$A[3] = 18$

consolidate: example




$A[0] = 7$

$A[1] = nil$

$A[2] = 17$

$A[3] = 18$

consolidate: example




$A[0] = nil$

$A[1] = (7)$

$A[2] = (17)$

$A[3] = (18)$

consolidate: example



$A[0] = \text{nil}$

$A[1] = (7)$

$A[2] = (17)$

$A[3] = (18)$

consolidate: algorithm

consolidate(H):


0. for each node n in H.root_list:
 1. x := n
 2. while A[x.degree] != null:
 3. y := A[x.degree]
 4. A[x.degree] := null
 5. if x.key > y.key:
 6. x, y := y, x
 7. remove y from H.root_list
 8. make y child of x // x.degree increases
 9. y.marked := false // used later
 10. A[x.degree] := x
 11. update H.min

decrease-priority: idea

- this is where we use the `marked` field
- `marked` is *true* if this node lost a child since being removed from root list
- cut child from parent: move child to root list and unmark it
- cascading cuts from some child node:
 - keep going up to root
 - if see an unmarked child, mark it and stop
 - if see a marked child, cut it and keep going

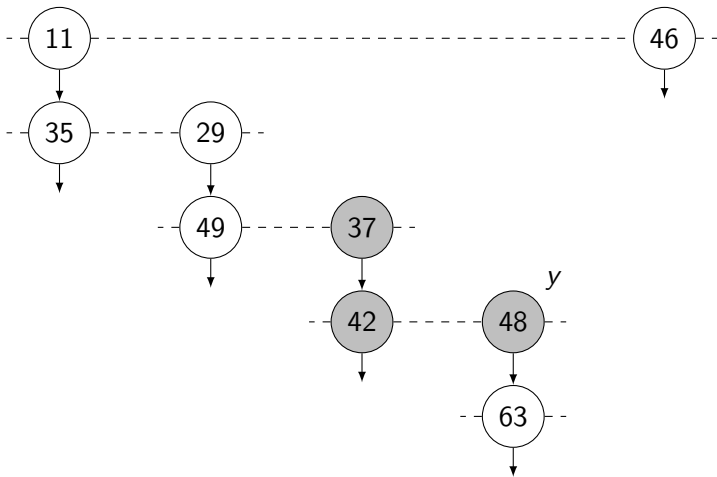
decrease-priority: example

decrease-priority(x , 46). $y.key > x.key$, will promote x .




decrease-priority: example

y lost a child while already marked, will be promoted.




decrease-priority: example

2nd y lost a child while already marked, will be promoted.



decrease-priority: example

3rd y lost a child while unmarked. Mark it now. Exit.



decrease-priority: algorithm

```
decrease-priority(H, x, k):  
0. if k >= x.key: return  
1. x.key := k  
2. y := x.parent  
3. if y != null and y.key > x.key:  
4.   cut(H, x, y)  
5.   while y.parent != null:  
6.     if not y.marked:  
7.       y.marked := true  
8.       break  
9.     else:  
10.      cut(H, y, y.parent)  
11.      y := y.parent  
12. if x.key < H.min.key:  
13.   H.min := x
```

decrease-priority: cut

```
cut(H, x, y):
```

0. remove x from children of y
1. add x to H.root_list
2. x.marked := false
3. if x.key < H.min.key:
4. H.min := x

complexity of Fibonacci heap operations

- Look at actual worst case time first
- Then define our potential function
- Then find amortised complexity of operations

complexity: actual costs

- define
 - $t(H)$: number of trees in heap H (nodes in the root list)
 - $d(H)$: degree of node with maximum degree in heap H
 - $D(n)$: maximum possible degree of Fibonacci Heap with n nodes
- $\text{insert}(j, p)$: $\mathcal{O}(1)$
- $\text{min}()$: $\mathcal{O}(1)$

complexity: actual costs

extract-min():

- remove node from root list: $\mathcal{O}(1)$
- insert children into root list: $d(H) \in \mathcal{O}(D(n))$
- now root list has: $t(H) + d(H)$ nodes
- consolidate(H):
 - how many times can a root become a child of another root? 1
 - max number of adoptions: $\mathcal{O}(t(H) + D(n))$
- H' : heap after running consolidate
- now root list has at most: $d(H') \in \mathcal{O}(D(n))$ nodes
- update $H'.min$: $\mathcal{O}(D(n))$
- total: $\mathcal{O}(t(H) + D(n))$

complexity: actual costs

define

- $m(H)$: number of marked nodes in heap

decrease-priority(n , p):

- set new priority of n : $\mathcal{O}(1)$
- if heap not ordered, cut n : $\mathcal{O}(1)$
- if cascading cuts: $\mathcal{O}(\#cuts)$
- only cut marked nodes during cascading cuts: $\#cuts \leq m(H)$
- so decrease-priority is $\mathcal{O}(m(H))$

complexity: observations

Observations AKA potential function magic:

- extract-min moves nodes from root list down
- decrease-priority cuts nodes / moves them up to root list
- extract-min: $\mathcal{O}(t(H) + D(n))$
- decrease-priority: $\mathcal{O}(m(H))$
- define potential function:

$$\Phi(H) = t(H) + 2 * m(H)$$

- Initially: $\Phi(H_0) = t(H_0) + 2 * m(H_0) = 0$

complexity: insert

- Potential function: $\Phi(H) = t(H) + 2 * m(H)$
- How does insert change potential:

$$\begin{aligned}\Delta(\Phi) &= \Phi(H_i) - \Phi(H_{i-1}) \\ &= t(H_i) + 2 * m(H_i) - t(H_{i-1}) - 2 * m(H_{i-1}) \\ &= 1\end{aligned}$$

- Then amortised complexity is:

$$a_i = t_i + \Delta(\Phi) = \mathcal{O}(1)$$

complexity: decrease-priority

- Potential function: $\Phi(H) = t(H) + 2 * m(H)$
- How does decrease-priority change potential?
- if we make x cuts
- for each cut, a node added to root list:

$$t(H_i) = t(H_{i-1}) + x$$

- every cut unmarks a marked node
- $x - 1$ or x nodes become unmarked
- at most 1 node becomes marked
- then:

$$m(H_i) \leq m(H_{i-1}) - (x - 1) + 1 = m(H_{i-1}) - x + 2$$

complexity: decrease-priority

Have:

- $\Phi(H) = t(H) + 2 * m(H)$
- $t(H_i) = t(H_{i-1}) + x$
- $m(H_i) \leq m(H_{i-1}) - x + 2$

Then:

$$\begin{aligned}\Delta(\Phi) &= \Phi(H_i) - \Phi(H_{i-1}) \\ &= t(H_i) + 2 * m(H_i) - t(H_{i-1}) - 2 * m(H_{i-1}) \\ &\leq t(H_{i-1}) + x + 2 * (m(H_{i-1}) - x + 2) - t(H_{i-1}) - 2 * m(H_{i-1}) \\ &= 4 - x\end{aligned}$$

Amortised cost:

$$a_i = t_i + \Phi(H_i) - \Phi(H_{i-1}) = (x + 1) + 4 - x = 5 \in \mathcal{O}(1)$$

complexity: extract-min

- Potential function: $\Phi(H) = t(H) + 2 * m(H)$
- How does extract-min change potential?
- no nodes become marked, some may become unmarked
 - $m(H_i) \leq m(H_{i-1})$
- after extract-min (after consolidate), all nodes in root list have different degree
- then: $t(H_i) \leq d(H_i) + 1$

complexity: extract-min

Then

$$\begin{aligned} a_i &= t_i + \Phi(H_i) - \Phi(H_{i-1}) \\ &\leq (t(H_{i-1}) + D(n)) + (t(H_i) + 2m(H_i)) - (t(H_{i-1}) + 2m(H_{i-1})) \\ &\leq t(H_{i-1}) + D(n) + D(n) + 1 + 2m(H_{i-1}) - t(H_{i-1}) - 2m(H_{i-1}) \\ &\in \mathcal{O}(D(n)) \end{aligned}$$

complexity: extract-min

- $a_i \in \mathcal{O}(D(n))$
- Last piece of the puzzle: a bound on $D(n)$
- What is the maximum degree of a root node in a heap of size n ?
- What is the minimum number of nodes $N(d)$ in a heap with root nodes of degree d ?
- ...
- In tutorial show $N(d) = \text{fib}(d + 2)$ — hence the name “Fibonacci heap”!
 $\therefore n \geq \phi^d$
 $\therefore d \leq \log_{\phi} n$

Fibonacci heap: complexity

- insert: amortised $\mathcal{O}(1)$
- extract-min: amortised $\mathcal{O}(\log n)$
- decrease-priority: amortised cost $\mathcal{O}(1)$